

A Problem-Oriented Perspective and Anchor Verification for Code Optimization

Tong Ye¹, Tengfei Ma², Xuhong Zhang¹, Hang Yu³,
Jianwei Yin¹, Wenhai Wang¹

¹Zhejiang University;

²Stony Brook University

³Ant Group

{tongye, zhangxuhong, zdzzlab}@zju.edu.cn, zjuyjw@cs.zju.edu.cn
tengfei.ma@stonybrook.edu, hyu.hugo@antgroup.com

Abstract

Large language models (LLMs) have shown remarkable capabilities in solving various programming tasks, such as code generation. However, their potential for code optimization, particularly in performance enhancement, remains largely unexplored. This paper investigates the capabilities of LLMs in optimizing code for minimal execution time, addressing a critical gap in current research. The recently proposed code optimization dataset constructs program optimization pairs based on iterative submissions from the same programmer for the same problem. However, this approach limits LLMs to local performance improvements, neglecting global algorithmic innovation. To overcome this limitation, we adopt a completely different perspective by reconstructing the optimization pairs into a problem-oriented approach. This allows for the integration of various ideas from multiple programmers tackling the same problem. Experimental results demonstrate that adapting LLMs to problem-oriented optimization pairs significantly enhances their optimization capabilities. Furthermore, recognizing the inherent trade-offs in code optimization, we introduce an anchor verification mechanism to mitigate the "optimization tax". Ultimately, our approach elevates both the optimization ratio and speedup to new levels.

1 Introduction

Large Language Models (LLMs) and Code LLMs, such as GPT-4 (Achiam et al., 2023), CodeLLaMA (Roziere et al., 2023), WizardCoder (Luo et al., 2024), DeepSeek-Coder (Guo et al., 2024) and Qwen2.5-Coder (Hui et al., 2024), have demonstrated remarkable capabilities in software engineering and programming tasks, garnering significant attention from both academia and industry. In tasks such as code completion and code generation, Code LLMs achieve high correctness rates (Pass@K) on widely used benchmarks like

EvalPlus (Liu et al., 2023) and LiveCodeBench (Jain et al., 2024). However, despite these advancements, the code produced by these models often falls short in real-world applications. It may lack the necessary optimizations to meet specific performance and efficiency requirements (Shi et al., 2024; Niu et al., 2024). As a result, the generated code often requires further refinement and optimization to align with practical constraints.

While low-level optimizing compilers and performance engineering tools have made significant advancements (Alfred et al., 2007; Wang and O’Boyle, 2018), they primarily focus on hardware-centric optimizations. High-level performance considerations, such as algorithm selection and API usage, continue to rely heavily on manual intervention by programmers. Automating high-level code optimization remains a major challenge and, unlike code generation, has yet to be widely explored. Code optimization can be approached from various angles. In this work, we specifically focus on time performance optimization, with an emphasis on minimizing program execution time, given its critical importance in practical applications.

In the field of performance optimization, the construction of datasets has been a critical challenge. Unlike code generation, which only requires the collection of correct code, performance optimization demands semantically equivalent code pairs with varying levels of efficiency. This dual requirement, ensuring both functional correctness and measurable performance improvements, makes dataset creation considerably more complex. Recent study (Shypula et al., 2024) partly addressed this challenge by collecting user iterative submissions from programming platforms, thereby creating code optimization pairs—each consisting of less efficient code and its semantically equivalent, more efficient counterpart. By utilizing these optimization pairs, researchers have demonstrated the potential of Code LLMs in code optimization tasks

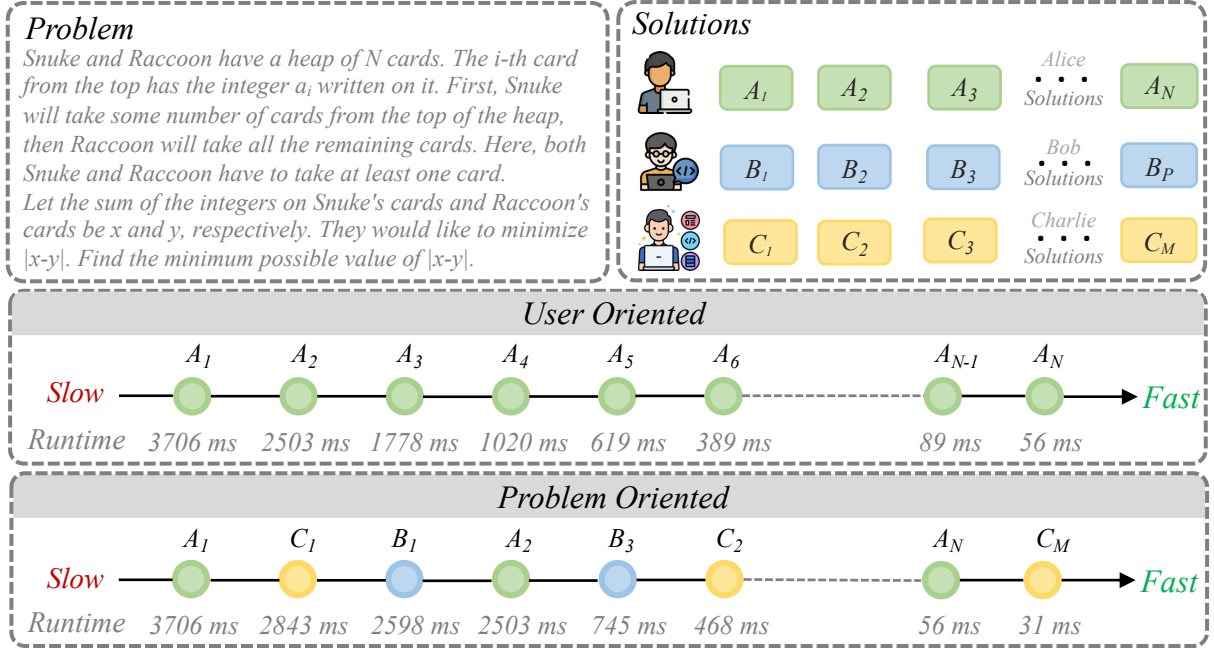


Figure 1: For a given problem, different users submit and iterate on their code solutions. The user-oriented perspective constructs optimization pairs based on the submission trajectories of individual users. In contrast, the problem-oriented perspective analyzes all solutions for the problem to build trajectories and form optimization pairs.

through subsequent fine-tuning.

However, the current approach of constructing code optimization pairs from iterative submissions by the same user has significant limitations. We refer to this as the **user-oriented** approach. As shown in Figure 1, a user initially submits a solution to a programming problem, but early versions may fail to meet the system’s time constraints due to excessive computational overhead. Through iterative refinements, the user eventually arrives at a more efficient solution. This process captures the user’s submission trajectory, which is used to construct optimization pairs such as (A_1, A_2) , (A_2, A_3) , ..., (A_{N-1}, A_N) . While this approach naturally reflects the direction of code optimization, it is inherently constrained by the thought patterns of a single programmer. Consequently, improvements tend to be incremental, building upon existing logic and paradigms. We present a substantial number of intuitive examples (Fig 15, 16, 17, 18) in Appendix H. In contrast, real-world code optimization thrives on collaborative diversity. Code review and refactoring processes deliberately involve multiple programmers to overcome cognitive inertia, with innovation arising from the synthesis of diverse perspectives. Inspired by this insight, we propose shifting from the user-oriented perspective to a **problem-oriented** perspective. We restructure optimization pairs by

incorporating solutions from multiple programmers addressing the same problem. As illustrated in the final part of Figure 1, solutions from different users, ordered by runtime, form a completely new optimization trajectory for the given problem. This problem-oriented perspective encourages a diverse range of innovative ideas, fostering a more holistic optimization process that better mirrors the complexity and creativity of program optimization.

Experimental results show that adapting Code LLMs to problem-oriented optimization pairs greatly enhances code optimization capabilities, leading to significant improvements in both optimization ratios ($31.24\% \rightarrow 58.89\%$) and speedup ($2.95\times \rightarrow 5.22\times$). Meanwhile, we also observe that code optimization inherently involves a trade-off in correctness, meaning that optimized code is not always guaranteed to be correct, which we call "optimization tax". To address this challenge, we introduce an innovative anchor verification mechanism specifically designed for code optimization. Unlike conventional test case execution feedback-based code generation methods (Chen et al., 2023; Wei et al., 2024; Chen et al., 2024a), which rely on synthesized test cases and bidirectional execution filtering to validate generated test case and code, our anchor verification mechanism first utilizes the LLM to explain the "slow code" and generate test case inputs. Next, we treat the "slow

code" under optimization as a test case anchor to produce precise outputs for these inputs. By pairing each test input with its corresponding output, we create complete and verified test cases. These verified test cases are then used for the iterative refinement of the "optimized code". Further experimental results show that our anchor verification mechanism pushes code optimization to new levels, significantly improving both the optimization ratio ($58.90\% \rightarrow 71.06\%$), speedup ($5.22\times \rightarrow 6.08\times$) and correctness ($61.55\% \rightarrow 74.54\%$).

In summary, our contributions are as follows, and the code is publicly available ¹:

- To the best of our knowledge, we are the first to introduce a problem-oriented perspective for code optimization.
- Our proposed anchor verification mechanism effectively mitigates the "optimization tax".
- Extensive experiments and analyses validate the effectiveness and robustness of both the problem-oriented perspective and the anchor verification mechanism in code optimization.

2 Related Works

2.1 LLMs for Code-Related Tasks.

LLMs pre-trained on extensive code corpora have demonstrated remarkable capabilities in various programming tasks, including code completion, code generation, and code summarization (Li et al., 2022; Nijkamp et al., 2023; Roziere et al., 2023; Wei et al., 2023; Guo et al., 2024; Song et al., 2024; Wang et al., 2025). To enhance the accuracy of code generation, numerous techniques and frameworks have been proposed, such as execution feedback and self-correction mechanisms (Chen et al., 2024b; Zhong et al., 2024; Moon et al., 2024; Olausson et al., 2024). However, despite these advancements, the research of LLMs to code optimization, a field of both practical significance and considerable real-world challenges, remains underexplored in both academia and industry.

2.2 Code Optimization.

With Moore’s law losing momentum, program optimization has become a central focus of software engineering over the past few decades (Bacon et al., 1994; Kistler and Franz, 2003). However, achieving high-level optimizations, such as algorithmic changes, remains challenging due to the difficulty

in comprehending code semantics. Previous research has employed machine learning to enhance performance by identifying compiler transformations (Bacon et al., 1994), optimizing GPU code (Liou et al., 2020), and automatically selecting algorithms (Kerschke et al., 2019). For instance, DeepPERF (Garg et al., 2022) leverages a transformer-based model fine-tuned to generate performance improvement patches for C# applications. Recently, Shypula et al. (2024) introduced the first C/C++ dataset designed for program efficiency optimization, with preliminary results demonstrating the potential of LLMs in code optimization.

3 Problem-Oriented Code Optimization

We shift from a user-oriented perspective to a problem-oriented perspective, with Section 3.1 outlining the key distinctions in their construction. Subsequently, we conduct comprehensive structural, semantic, and sampling analyses of both user-oriented and problem-oriented optimization pairs.

3.1 Problem-oriented Optimization Pairs

User-Oriented Perspective. The current code optimization pairs are derived from PIE, introduced by Shypula et al. (2024), which focuses on optimizing program execution time by utilizing human programmers’ submissions from a wide range of competitive programming tasks on CodeNet (Puri et al., 2021). A key aspect of developing PIE is recognizing the typical workflow of programmers: when faced with a problem, they usually begin with an initial solution and then iteratively refine it. As shown in Figure 1, for a given problem \mathcal{P} , users (Alice, Bob, Charlie, etc.) have their submission trajectories, filter out incorrect submissions, and sort the remaining ones in chronological order. Formally denoted as:

Alice valid submissions: $[A_1, A_2, A_3, \dots, A_N]$

Bob valid submissions: $[B_1, B_2, B_3, \dots, B_P]$

Charlie valid submissions: $[C_1, C_2, C_3, \dots, C_M]$

The user-oriented optimization pairs are constructed by extracting sequential pairs from each user’s submission trajectory. For example, Alice’s valid submissions generate optimization pairs such as (A_1, A_2) , (A_2, A_3) , and so on, while Charlie’s valid submissions result in optimization pairs like (C_1, C_2) , (C_2, C_3) , and so forth. Ultimately, aggregating all these optimization pairs forms the complete user-oriented optimization dataset (PIE).

¹<https://anonymous.4open.science/r/code-optimization-85ED>




Figure 2: Structural Analysis of the Disparities between Problem-oriented and User-oriented Optimization Pairs.

Problem-Oriented Perspective. While user-oriented optimization pairs naturally indicate the direction of optimization, as previously noted, they are inherently confined by the cognitive patterns of a single programmer. The detailed instances in Appendix H illustrate this point, intuitively showing that the overall problem-solving approach and logical framework remain largely unaltered. Therefore, we shift the perspective on optimization pairs and propose a problem-oriented construction method. Specifically, we regard all submissions for the same problem \mathcal{P} from different users as a single group, thereby breaking down the barriers between different users. We sort all valid user submissions for the same \mathcal{P} based on the marked runtime and map them onto the same optimization trajectories:

$$\text{All users: } [A_1, C_1, B_1, A_2, B_3, C_2, \dots, C_M]$$

Subsequently, we construct optimization pairs along the problem-oriented trajectory, such as $(A_1, C_1), (C_1, B_1), (C_1, B_2), \text{etc.}$ Ultimately, this process yields a problem-oriented optimization dataset. This method not only reflects the direction of optimization but also integrates the diverse strategies and algorithms of different programmers.

Extra Quantity Bonus. The problem-oriented perspective also offers a significant advantage in terms of scale. Assuming there are \mathcal{P} problems, each with \mathcal{U} users, and each user has n_u valid submissions, the user-oriented and problem-oriented perspectives exhibit a substantial divergence in the scaling of optimization pairs:

$$\begin{aligned} \# \text{ user oriented} &= \frac{1}{2} \cdot \sum_{p=1}^{\mathcal{P}} \sum_{u=1}^{\mathcal{U}} C_{n_u}^2 \\ \# \text{ problem oriented} &= \frac{1}{2} \cdot \sum_{p=1}^{\mathcal{P}} C_{\sum_{u=1}^{\mathcal{U}} n_u}^2 \end{aligned}$$




Figure 3: Semantic Representation Analysis of Problem-oriented and User-oriented Optimization Pairs.




Figure 4: Human Analysis of the Optimization Types between Problem-oriented and User-oriented Pairs.

This characteristic is particularly advantageous for addressing the challenge of data scarcity when constructing code optimization pairs.

3.2 Multi-Dimension Analysis

To rigorously and comprehensively compare code optimization pairs derived from two different perspectives, we employ a multi-faceted analysis. Specifically, based on the problem-oriented approach proposed in Section 3.1, we reconstruct the PIE dataset, resulting in the PCO (Problem-oriented Code Optimization). To ensure comparability and fairness, we retained the same number of optimization pairs for each problem in PCO as in the corresponding problem in PIE, selecting those with the top speedup rankings. This guarantees that both datasets contain a total of 78K optimization pairs, as shown in Table 3. We then perform comparative analyses across three different dimensions: Structural Analysis, Semantic Representation Analysis, and Human & LLMs Sampling Analysis.

Table 1: Prompt and Fine-Tuning Results for LLMs on PIE and PCO optimization pairs with BEST@1 and BEST@8.

Prompt / Dataset	LLMs & Code LLMs	BEST@1			BEST@8		
		%OPT	SPEEDUP	CORRECT	%OPT	SPEEDUP	CORRECT
Instruct	DEEPSEEKCODER 33B	5.28%	1.12×	30.17%	14.83%	1.23×	48.00%
Instruct	GPT-4	12.37%	1.19×	75.28%	22.81%	1.38×	91.74%
CoT	DEEPSEEKCODER 33B	13.91%	1.24×	37.45%	20.81%	1.55×	61.89%
CoT	GPT-4	23.43%	1.37×	48.65%	47.92%	1.74×	80.53%
PIE	CODELLAMA 13B	12.98%	1.73×	47.45%	41.65%	2.85×	72.27%
PIE	DEEPSEEKCODER 7B	23.56%	2.29×	41.27%	47.23%	3.34×	69.23%
PIE	DEEPSEEKCODER 33B	27.57%	2.77×	50.49%	56.76%	3.83×	81.14%
PIE	QWEN2.5-CODER 7B	26.96%	2.80×	41.21%	56.17%	3.85×	78.54%
PIE	QWEN2.5-CODER 32B	31.24%	2.95×	46.52%	60.89%	4.11×	87.95%
PCO	CODELLAMA 13B	31.83%	3.23×	44.26%	55.87%	4.89×	69.61%
PCO	DEEPSEEKCODER 7B	44.38%	4.31×	45.71%	71.53%	6.24×	73.09%
PCO	DEEPSEEKCODER 33B	49.83%	4.57×	50.64%	74.87%	6.67×	78.29%
PCO	QWEN2.5-CODER 7B	54.83%	4.73×	56.26%	75.28%	6.89×	77.43%
PCO	QWEN2.5-CODER 32B	58.89%	5.22×	61.55%	80.77%	7.22×	83.03%

Structural Analysis. First, we analyze the structural differences between "slow" and "fast" code within the optimization pairs. To achieve this, we utilize Control Flow Graphs (CFGs), as they represent the logical structure and execution pathways of a program. To quantify the structural differences, we employ the Graph Edit Distance (GED) metric, which measures the minimum edit operation cost between the CFGs of "slow" and "fast" code. As shown in Fig 2, significant differences emerge from different perspectives: user-oriented optimization pairs exhibit a relatively small average GED, indicating that the optimizations involve minor changes, such as localized optimizations. In contrast, problem-oriented optimization pairs show a significantly higher average GED, suggesting substantial changes, such as major structural modifications. These optimizations often involve global changes, such as algorithmic adjustments, which contrasts sharply with the incremental nature of user-oriented optimizations.

Semantic Representation Analysis. Beyond comparing the code structure within optimization pairs, we also analyze the semantic differences between optimization pairs. To do this, we concatenate the "slow" and "fast" code snippets within each pair to form a unified input. These concatenated sequences are then encoded using CODET5P-110M-EMBEDDING (Wang et al., 2023) to generate high-dimensional semantic embeddings, which are subsequently projected using t-SNE (van der Maaten and Hinton, 2008) for visualization. As shown in Fig 3, the embeddings for the user-oriented pairs

are tightly clustered, indicating that the code pairs represent similar coding semantics. In contrast, the embeddings for the problem-oriented pairs are more dispersed, reflecting greater diversity.

Human & LLMs Sampling Analysis. We conduct a sampling analysis to further investigate optimization patterns. Specifically, we randomly select 100 pairs from the PIE and PCO for human analysis, aiming to classify the types of optimizations applied. Additionally, we randomly select 1,000 optimization pairs for evaluation using GPT-4. The optimizations are categorized into three main types: global algorithmic optimizations, local optimizations, and other modifications (e.g., code cleanup), with details provided in the Appendix A. As shown in Fig 4, human analysis reveals distinct trends across the different perspectives: In the PIE dataset, true global algorithmic optimizations constitute a relatively small proportion. In contrast, the majority of program pairs in PCO fall into the global algorithmic optimization category, indicating a stronger emphasis on significant algorithmic and structural improvements. The LLM analysis exhibits similar patterns, as shown in Fig 7.

3.3 Adapting LLMs to Optimization Pairs

Moreover, we utilize supervised finetuning to adapt LLMs to problem-oriented PCO optimization pairs.

Metrics. To evaluate code optimization performance, following (Shypula et al., 2024), we measure below metrics:

- **Percent Optimized [%OPT]:** The fraction of programs in the test set improved by a certain




Figure 5: Impact of using varying percentages of PCO optimization pairs on %OPT, SPEEDUP, and CORRECT. The blue line represents the original PCO datasets, while the yellow line represents the original PIE datasets.

method. A program must be at least 10% faster and correct to contribute.

- **Speedup** [SPEEDUP]: The absolute improvement in running time. If o and n are the "old" and "new" running times, then $\text{SPEEDUP}(o, n) = \left(\frac{o}{n}\right)$. A program must be correct to contribute.
- **Percent Correct** [CORRECT]: The proportion of programs in the test set that are functionally equivalent to the original program.

We count a program as functionally correct only if it passes every test case. Additionally, we report SPEEDUP as the average speedup across all test set samples. For generated programs that are either incorrect or slower than the original, we use a speedup of $1.0\times$, hence, in the worst case, the original program has a speedup of 1.0 . We benchmark performance using gem5 CPU simulator environment (Binkert et al., 2011) and compile all C++ programs with GCC version 9.4.0 and C++17 as well as the $-O3$ optimization flag. Therefore, any reported improvements would be those on top of the optimizing compiler.

Code LLMs Selection. We select GPT-4 (*gpt-4-0613*) (Achiam et al., 2023), CODELLAMA (Roziere et al., 2023), DEEPSEEKCODER (Guo et al., 2024) and QWEN2.5-CODER (Hui et al., 2024) for code optimization, as these LLMs are top-performing in code domain. For instruction-following prompt, we utilize the corresponding *chat* versions, while for fine-tuning, we employ the *base* versions of these LLMs. Detailed training parameters are provided in the Appendix D.

Decoding Strategy. Code generation benefits from sampling multiple candidate outputs for each input and selecting the best one; in our case, the "best" refers to the fastest program that passes all

test cases. We use $\text{BEST}@k$ to denote this strategy, where k represents the number of samples and the temperature is set to 0.7. we use vLLM (Kwon et al., 2023) for efficiently inference and detailed prompts shown in Fig 10.

3.4 Adapting Results.

Instruction Prompting. First, we use instruction prompts to guide the LLMs in optimizing code. Additionally, inspired by Chain-of-Thought (CoT) (Wei et al., 2022), we ask the LLMs to reason about how to optimize the program before generating the optimized version. Details of the instruction prompt and CoT prompt are shown in Appendix E. Table 1 shows that using instruct prompt and CoT did not significantly improve %OPT and SPEEDUP. The best performance by GPT-4 achieved 47.92 %OPT and $1.74\times$ SPEEDUP under $\text{BEST}@8$. Additionally, we observe that using CoT for optimization speeds up the program but can lead to a decline in CORRECT due to the complexities it introduces.

Fine-Tuning Results. As shown in Table 1, whether for different LLM series or varying parameter scales, significant performance differences are observed when finetuned on user-oriented (PIE) and problem-oriented (PCO) optimization pairs. QWEN2.5-CODER 32B on PCO at $\text{BEST}@1$, demonstrates substantial improvements: %OPT ($31.24\% \rightarrow 58.89\%$), SPEEDUP ($2.95\times \rightarrow 5.22\times$), and CORRECT ($46.52\% \rightarrow 61.55\%$) compared to finetuned on PIE. At $\text{BEST}@8$, %OPT and SPEEDUP reached 80.77% and $7.22\times$, respectively. This indicates a significant advantage in adapting to problem-oriented optimization pairs compared to user-oriented optimization pairs.

Finding 1: We also observe that, unlike $\text{BEST}@1$, CORRECT slightly declines for most LLMs adapted on PCO under $\text{BEST}@8$. This is because, com-

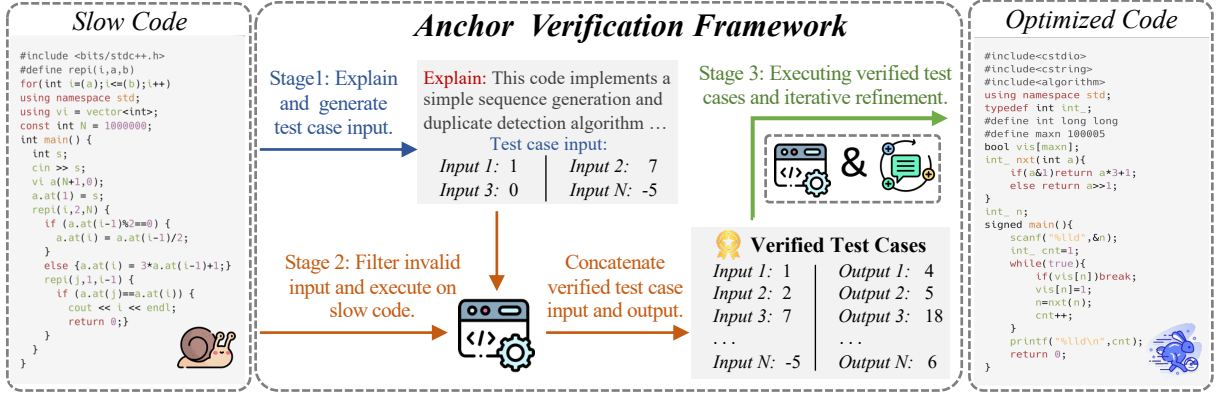


Figure 6: Anchor Verification Framework. It includes three stages: generating test inputs based on the slow code’s functionality, constructing a verified test case set by executing inputs through the slow code, and iteratively refining the optimized code with execution feedback to ensure correctness and preserve performance gains.

pared to PIE, LLMs adapting on PCO results in more significant modifications to the code in pursuit of maximum efficiency, which slightly disrupts the balance of CORRECT.

Finding 2: For LLMs adapted on PCO, both %OPT and CORRECT are much closer compared to PIE. This suggests that when the optimized code is correct, it is highly likely to be optimized. The closer %OPT and CORRECT are, the higher the proportion of "correct will be optimized". This insight also indicates that, for LLMs adapted on PCO, to further increase the optimization ratio, the bottleneck lies in ensuring correctness.

3.5 PCO Percentage Analysis.

We further explore how using fewer PCO optimization pairs impacts %OPT, SPEEDUP, and CORRECT. To investigate this, we randomly selected a certain percentage of optimization pairs from PCO, reducing the number of pairs from 90% and 80% down to 10%, and fine-tuned QWEN2.5-CODER in the same way. As shown in Fig 5, even with just 30% of the PCO optimization pairs, LLMs adapted on PCO achieve both %OPT and SPEEDUP that surpass those of the full PIE. Furthermore, with roughly half of the PCO pairs, CORRECT matches the full PIE. These results highlight the impressive data efficiency of the problem-oriented perspective, where fewer optimization pairs can still deliver competitive or even superior performance compared to full user-oriented optimization pairs.

4 Anchor Verification For Practicability

In Section 3, we discuss a key challenge in LLM code optimization, whether through Instruct Prompting or Finetuning: while performance en-

hancement offers significant benefits, there is also a risk of optimized code is not 100% correct. We refer to this phenomenon as the "optimization tax". To tackle the challenge of "optimization tax", we propose a novel anchor verification framework that leverages the original "slow code" as a reliable test case verification anchor. Unlike the code generation domain, which may rely on potentially error-prone synthetic test cases for refinement, the code optimization scenario has a unique advantage: the "slow code", despite its inefficiency, is functionally correct. This inherent characteristic positions it as an ideal test case verification anchor. As shown in Fig 6, the framework consists of three stages. First, we generate test case inputs (only inputs) based on the functionality of the slow code. Then, we construct a verified test case (test case inputs and outputs) by executing these inputs through the slow code. Finally, we iteratively refine the optimized code using feedback from the execution of the verified test case. This process ensures correctness while maintaining the performance gains.

4.1 Detailed Methodology.

Stage 1: Test Inputs Generation. In the first stage, the LLM is prompted to explain the functionality of the "slow code" and generate a set of test inputs. These test inputs are designed to cover the boundary cases of the implemented functionality. Unlike LLM-based test case generation, the first stage focuses solely on generating the test case inputs.

Stage 2: Verified Testcase Construction. Based on the obtained test case inputs, we feed these inputs to the "slow code" for compilation and execution. Although the "slow code" is inefficient, it ensures correctness. We filter out test case inputs

Table 2: Results of Anchor Verification and compared methods with QWEN2.5-CODER, GPT-4o, and DEEPSEEK-V3 on BEST@1. The improvement (denoted as Δ) is measured against the baseline (w/o refinement).

LLMs	Methods	BEST@1					
		%OPT	$\Delta \uparrow$	SPEEDUP	$\Delta \uparrow$	CORRECT	$\Delta \uparrow$
	Baseline (w/o refinement)	58.90%		5.22×		61.55%	
QWEN2.5-CODER 32B INSTRUCT	Self Debugging	58.42%	-0.48	5.13×	-0.09	61.14%	-0.41
	Direct Test Generation	62.98%	+4.08	5.46×	+0.24	65.95%	+4.40
	Anchor Verification	64.75%	+5.85	5.67×	+0.45	67.28%	+5.73
GPT-4o	Self Debugging	61.96%	+3.06	5.59×	+0.37	63.60%	+2.05
	Direct Test Generation	65.43%	+6.53	5.71×	+0.49	68.61%	+7.06
	Anchor Verification	68.40%	+9.50	5.90×	+0.68	71.98%	+10.43
DEEPSEEK-V3	Self Debugging	64.11%	+5.21	5.63×	+0.41	65.64%	+4.09
	Direct Test Generation	66.26%	+7.36	5.81×	+0.59	69.53%	+7.98
	Anchor Verification	71.06%	+12.16	6.08×	+0.86	74.54%	+12.99

that don't match the input format and gather the corresponding output results. Finally, we combine the test case inputs and their corresponding output results to form verified test case sets.

Stage 3: Iterative Refinement. Leveraging the verified test case sets, we compile and execute the optimized code to check its correctness. If an error occurs, similar to the feedback mechanism in code generation, we provide the error information to the LLM, enabling it to iteratively refine the optimized code based on this feedback.

4.2 Experiment Results.

Compared Methods. To rigorously validate the effectiveness of the anchor verification mechanism, we benchmark against two compared methods:

- **Self-Debugging:** following the approach outlined in (Chen et al., 2024b), the method prompts the LLM to provide line-by-line explanations of the generated code as feedback for refinement.
- **Direct Test Generation:** The LLM generates complete test cases (including inputs and outputs) and uses synthetic cases to execute and iteratively refine the optimized code.

Experiments Setup. We set the maximum iteration count for all methods to 1. Detailed implementation and prompts are shown in Appendix F.

Main Results. We use "QWEN2.5-CODER 32B finetuned on PCO" as the baseline (the last row in Table 1). We experimented with three different LLM backbones: QWEN2.5-CODER 32B INSTRUCT, GPT-4o, and DEEPSEEK-V3 (DeepSeek-AI, 2024), with the results shown in Table 2. All methods showed performance gains, except for a slight decline in the self-debugging with QWEN2.5-

CODER 32B INSTRUCT. The decline can be attributed to the high demands on the LLM's ability for self-explanation and correction, and QWEN2.5-CODER 32B INSTRUCT's overall performance still lags behind the other two LLMs. Anchor verification demonstrated the best improvements across all three LLM backbones, particularly with DEEPSEEK-V3. Compared to the baseline, CORRECT improved by 12.99%, %OPT improved by 12.16%, and SPEEDUP increased to 6.08×. This result further confirms that improving CORRECT can simultaneously enhance both %OPT and SPEEDUP.

4.3 Deeper Analysis.

To conduct a more comprehensive analysis of Anchor Verification, we perform experiments using "QWEN2.5-CODER 32B finetuned on PIE" as the baseline and compared it with other methods. The results, presented in Table 4, show that Anchor Verification continues to deliver the highest performance gains. Furthermore, we observed that the gains in optimization ratio and speedup brought by Anchor Verification's improvement in correctness in PIE are not as significant as those achieved in PCO. This further underscores the superiority of the PCO. Additionally, we present three case studies to intuitively show specific examples of Anchor Verification, as shown in Figure 19, 20, and 21.

5 Conclusion

In this paper, we propose a problem-oriented perspective and an anchor verification mechanism for code optimization. Our approach elevates both the optimization ratio, speedup, and correctness to new levels. We hope these insights will pave the way for a feasible path to improving program efficiency.

Limitation

This paper focuses on optimizing the time efficiency of given code, without considering other optimization directions. However, in real-world scenarios, there are many other optimization directions, such as memory optimization. Furthermore, ensuring the full correctness of code optimization remains a complex challenge, one that warrants further research.

References

- Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, et al. 2023. [Gpt-4 technical report](#). *arXiv preprint arXiv:2303.08774*.
- V Aho Alfred, S Lam Monica, and D Ullman Jeffrey. 2007. [Compilers Principles, Techniques & Tools](#). pearson Education.
- David F. Bacon, Susan L. Graham, and Oliver J. Sharp. 1994. [Compiler transformations for high-performance computing](#). *ACM Comput. Surv.*, 26(4):345–420.
- Nathan Binkert, Bradford Beckmann, Gabriel Black, Steven K. Reinhardt, Ali Saidi, Arkaprava Basu, Joel Hestness, Derek R. Hower, Tushar Krishna, Somayeh Sardashti, Rathijit Sen, Korey Sewell, Muhammad Shoaib, Nilay Vaish, Mark D. Hill, and David A. Wood. 2011. [The gem5 Simulator](#). *SIGARCH Comput. Archit. News*.
- Bei Chen, Fengji Zhang, Anh Nguyen, Daoguang Zan, Zeqi Lin, Jian-Guang Lou, and Weizhu Chen. 2023. [Codet: Code generation with generated tests](#). In *The Eleventh International Conference on Learning Representations*.
- Mouxian Chen, Zhongxin Liu, He Tao, Yusu Hong, David Lo, Xin Xia, and Jianling Sun. 2024a. [B4: Towards optimal assessment of plausible code solutions with plausible tests](#). In *Proceedings of the 39th IEEE/ACM International Conference on Automated Software Engineering, ASE '24*, page 1693–1705. ACM.
- Xinyun Chen, Maxwell Lin, Nathanael Schärli, and Denny Zhou. 2024b. [Teaching large language models to self-debug](#). In *The Twelfth International Conference on Learning Representations*.
- DeepSeek-AI. 2024. [Deepseek-v3 technical report](#). *Preprint*, arXiv:2412.19437.
- Spandan Garg, Roshanak Zilouchian Moghaddam, Colin B. Clement, Neel Sundaresan, and Chen Wu. 2022. [Deepperf: A deep learning-based approach for improving software performance](#). *Preprint*, arXiv:2206.13619.
- Daya Guo, Qihao Zhu, Dejian Yang, Zhenda Xie, Kai Dong, Wentao Zhang, Guanting Chen, Xiao Bi, Y Wu, YK Li, et al. 2024. [Deepseek-coder: When the large language model meets programming—the rise of code intelligence](#). *arXiv preprint arXiv:2401.14196*.
- Edward J Hu, yelong shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2022. [LoRA: Low-rank adaptation of large language models](#). In *International Conference on Learning Representations*.
- Binyuan Hui, Jian Yang, Zeyu Cui, Jiaxi Yang, Dayiheng Liu, Lei Zhang, Tianyu Liu, Jiajun Zhang, Bowen Yu, Kai Dang, et al. 2024. [Qwen2. 5-coder technical report](#). *arXiv preprint arXiv:2409.12186*.
- Naman Jain, King Han, Alex Gu, Wen-Ding Li, Fanjia Yan, Tianjun Zhang, Sida Wang, Armando Solar-Lezama, Koushik Sen, and Ion Stoica. 2024. [Livecodebench: Holistic and contamination free evaluation of large language models for code](#). *arXiv preprint arXiv:2403.07974*.
- Pascal Kerschke, Holger H. Hoos, Frank Neumann, and Heike Trautmann. 2019. [Automated algorithm selection: Survey and perspectives](#). *Evolutionary Computation*, 27(1):3–45.
- Thomas Kistler and Michael Franz. 2003. [Continuous program optimization: A case study](#). *ACM Trans. Program. Lang. Syst.*, 25(4):500–548.
- Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. 2023. [Efficient memory management for large language model serving with pagedattention](#). In *Proceedings of the ACM SIGOPS 29th Symposium on Operating Systems Principles*.
- Yujia Li, David Choi, Junyoung Chung, Nate Kushman, Julian Schrittwieser, Rémi Leblond, Tom Eccles, James Keeling, Felix Gimeno, Agustin Dal Lago, et al. 2022. [Competition-level code generation with alphacode](#). *Science*, 378(6624):1092–1097.
- Jhe-Yu Liou, Xiaodong Wang, Stephanie Forrest, and Carole-Jean Wu. 2020. [Gevo: Gpu code optimization using evolutionary computation](#). *ACM Trans. Archit. Code Optim.*, 17(4).
- Jiawei Liu, Chunqiu Steven Xia, Yuyao Wang, and Lingming Zhang. 2023. [Is your code generated by chat-GPT really correct? rigorous evaluation of large language models for code generation](#). In *Thirty-seventh Conference on Neural Information Processing Systems*.
- Ilya Loshchilov and Frank Hutter. 2019. [Decoupled weight decay regularization](#). In *International Conference on Learning Representations*.
- Ziyang Luo, Can Xu, Pu Zhao, Qingfeng Sun, Xiubo Geng, Wenxiang Hu, Chongyang Tao, Jing Ma, Qingwei Lin, and Daxin Jiang. 2024. [Wizardcoder](#):

- Empowering code large language models with evol-instruct. In *The Twelfth International Conference on Learning Representations*.
- Seungjun Moon, Hyungjoo Chae, Yongho Song, Taeyoon Kwon, Dongjin Kang, Kai Tzu iunn Ong, Seung won Hwang, and Jinyoung Yeo. 2024. [Coffee: Boost your code llms by fixing bugs with feedback](#). Preprint, arXiv:2311.07215.
- Erik Nijkamp, Bo Pang, Hiroaki Hayashi, Lifu Tu, Huan Wang, Yingbo Zhou, Silvio Savarese, and Caiming Xiong. 2023. [Codegen: An open large language model for code with multi-turn program synthesis](#). In *The Eleventh International Conference on Learning Representations*.
- Changan Niu, Ting Zhang, Chuanyi Li, Bin Luo, and Vincent Ng. 2024. [On evaluating the efficiency of source code generated by llms](#). Preprint, arXiv:2404.06041.
- Theo X. Olausson, Jeevana Priya Inala, Chenglong Wang, Jianfeng Gao, and Armando Solar-Lezama. 2024. [Is self-repair a silver bullet for code generation?](#) In *The Twelfth International Conference on Learning Representations*.
- Ruchir Puri, David S Kung, Geert Janssen, Wei Zhang, Giacomo Domeniconi, Vladimir Zolotov, Julian Dolby, Jie Chen, Mihir Choudhury, Lindsey Decker, et al. 2021. [Codenet: A large-scale ai for code dataset for learning a diversity of coding tasks](#). arXiv preprint arXiv:2105.12655.
- Baptiste Roziere, Jonas Gehring, Fabian Gloeckle, Sten Sootla, Itai Gat, Xiaoqing Ellen Tan, Yossi Adi, Jingyu Liu, Tal Remez, Jérémy Rapin, et al. 2023. [Code llama: Open foundation models for code](#). arXiv preprint arXiv:2308.12950.
- Jieke Shi, Zhou Yang, and David Lo. 2024. [Efficient and green large language models for software engineering: Literature review, vision, and the road ahead](#). Preprint, arXiv:2404.04566.
- Alexander G Shypula, Aman Madaan, Yimeng Zeng, Uri Alon, Jacob R. Gardner, Yiming Yang, Milad Hashemi, Graham Neubig, Parthasarathy Ranganathan, Osbert Bastani, and Amir Yazdanbakhsh. 2024. [Learning performance-improving code edits](#). In *The Twelfth International Conference on Learning Representations*.
- Zifan Song, Yudong Wang, Wenwei Zhang, Kuikun Liu, Chengqi Lyu, Demin Song, Qipeng Guo, Hang Yan, Dahua Lin, Kai Chen, and Cairong Zhao. 2024. [Alchemistcoder: Harmonizing and eliciting code capability by hindsight tuning on multi-source data](#). In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*.
- Laurens van der Maaten and Geoffrey Hinton. 2008. [Visualizing data using t-sne](#). *Journal of Machine Learning Research*, 9(86):2579–2605.
- Yaoxiang Wang, Haoling Li, Xin Zhang, Jie Wu, Xiao Liu, Wenxiang Hu, Zhongxin Guo, Yangyu Huang, Ying Xin, Yujiu Yang, Jinsong Su, Qi Chen, and Scarlett Li. 2025. [Epicoder: Encompassing diversity and complexity in code generation](#). Preprint, arXiv:2501.04694.
- Yue Wang, Hung Le, Akhilesh Gotmare, Nghi Bui, Junnan Li, and Steven Hoi. 2023. [CodeT5+: Open code large language models for code understanding and generation](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 1069–1088, Singapore. Association for Computational Linguistics.
- Zheng Wang and Michael O’Boyle. 2018. [Machine learning in compiler optimization](#). *Proceedings of the IEEE*, 106(11):1879–1901.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, brian ichter, Fei Xia, Ed Chi, Quoc V Le, and Denny Zhou. 2022. [Chain-of-thought prompting elicits reasoning in large language models](#). In *Advances in Neural Information Processing Systems*, volume 35, pages 24824–24837. Curran Associates, Inc.
- Yuxiang Wei, Federico Cassano, Jiawei Liu, Yifeng Ding, Naman Jain, Zachary Mueller, Harm de Vries, Leandro von Werra, Arjun Guha, and Lingming Zhang. 2024. [Selfcodealign: Self-alignment for code generation](#). arXiv preprint arXiv:2410.24198.
- Yuxiang Wei, Zhe Wang, Jiawei Liu, Yifeng Ding, and Lingming Zhang. 2023. [Magicoder: Source code is all you need](#). arXiv preprint arXiv:2312.02120.
- Yaowei Zheng, Richong Zhang, Junhao Zhang, Yanhan Ye, Zheyang Luo, Zhangchi Feng, and Yongqiang Ma. 2024. [Llamafactory: Unified efficient fine-tuning of 100+ language models](#). In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 3: System Demonstrations)*, Bangkok, Thailand. Association for Computational Linguistics.
- Li Zhong, Zilong Wang, and Jingbo Shang. 2024. [Ldb: A large language model debugger via verifying runtime execution step-by-step](#). arXiv preprint arXiv:2402.16906.

A Categories of Optimization Types.

We categorize code optimization into three main categories: global algorithmic optimizations, local optimizations, and other optimizations.

- **Global Algorithmic Optimizations:** This type of optimization involves altering the algorithm itself to achieve significant performance improvements. Such changes can effectively reduce time complexity and enhance the speed of code execution. Examples include transforming recursive solutions into dynamic programming ap-




Figure 7: LLM Analysis of the Optimization Types between Problem-oriented and User-oriented Pairs.

proaches, leveraging advanced mathematical theories, and restructuring complex data processing logic. These optimizations can lead to substantial gains in efficiency and scalability.

- **Local Optimizations:** These optimizations focus on improving specific parts of the code without changing the overall algorithm. They include enhancing I/O functions, optimizing read/write patterns to minimize runtime delays, and reducing computational complexity in certain sections of the code. By addressing these localized issues, programs can achieve more efficient execution and better resource utilization, ultimately leading to faster and more responsive applications.
- **Other Optimizations:** This category involves general code cleanup and refactoring aimed at improving code readability, maintainability, and overall quality. Examples include removing unnecessary initializations and redundant code, cleaning up outdated comments, and organizing the code structure more logically.

B LLMs Analysis on Optimization Types.

Figure 7 presents the LLMs analysis of optimization types between problem-oriented and user-oriented optimization pairs. GPT-4 identifies a higher proportion of "global algorithm optimization" compared to human analysis. Upon further investigation, we find that this discrepancy is mainly due to GPT-4's tendency to categorize program pairs with significant changes as "global algorithm optimization".

C Datasets Statistics.

The statistical results of the PCO and PIE are shown in Table 3. We meticulously reviewed and ensured

Given the program below, improve
→ its performance:

Program:
{slow_code}

Optimized Version:

Figure 8: Instruct Prompt.

that any particular competitive programming problem appeared in only one of the train, validation, or test sets.

Dataset	Unique Problems	Pairs
PIE	1,474	77,967
PCO	1,474	77,967
Val	77	2,544
Test	41	978

Table 3: Number of unique problem ids and pairs.

D Training Details.

We fine-tuned the CODELLAMA (13B), DEEPSEEK-CODER (7B, 33B), and QWEN2.5-CODER (7B, 32B) models using LLAMA-FACTORY (Zheng et al., 2024) on a server equipped with 8×A100 GPUs (NVIDIA A100 80GB). During the fine-tuning process, we employed LoRA (Hu et al., 2022) (with lora_rank=8 and lora_target=all), and for both the PIE and PCO datasets, we trained the LLMs for only 2 epochs. All experiments were conducted using AdamW (Loshchilov and Hutter, 2019) optimizer with an initial learning rate 5e-5.

E The Prompts of Adapting LLM on Optimization Pairs.

In this section, we present the prompts for adapting the LLM to optimization pairs. The instruction prompt is shown in Figure 8, the CoT (Chain of Thought) prompt is shown in Figure 9, and the vLLM inference prompt is shown in Figure 10.

F Implementation Details of the Anchor Verification Framework and the Compared Methods.

- **Anchor Verification:** In the Anchor Verification Framework, for the test case inputs in Stage 1,

Given the program, generate an
 ↳ efficiency improvement strategy
 ↳ to enhance its performance.

slower program:
 {slow_code}

strategy:
 LLMs generated potential strategy.

optimized version:

Figure 9: Chain-of-thought Prompting.

Given the program below, improve
 ↳ its performance:

Program:
 {slow_code}

Optimized Version:

Figure 10: Inference Prompt.

we prompt the LLM to generate three test case inputs based on the "slow code", the detailed prompt as illustrated in Figure 11. In Stage 2 and Stage 3, for compiling and executing both the "slow code" and "optimized code", we compile all C++ programs using GCC version 9.4.0 with C++17 and the -O3 optimization flag. In Stage 3, we leverage the verified test case sets. If an error occurs, we provide the error information to the LLM, allowing it to iteratively refine the optimized code based on this feedback. The detailed prompt is shown in Figure 12.

- **Self-Debugging:** following the approach presented in (Chen et al., 2024b), the method instructs the LLM to provide line-by-line explanations of the generated program as feedback, functioning akin to rubber duck debugging. In this process, the LLM is capable of autonomously identifying and rectifying bugs without requiring human intervention. The detailed prompt is shown in Figure 13.
- **Direct Test Generation:** The LLM generates complete test cases (including both inputs and outputs) and utilizes synthetic test cases to execute the optimized code, enabling iterative refinement. The prompt for generating complete test cases is shown in Figure 14, while the iterative refinement prompt is the same as the one used

Given the program below, please
 ↳ explain and analyze its
 ↳ functionality, and provide 3
 ↳ testcase inputs that fully
 ↳ consider boundary conditions
 ↳ and code coverage. Note that
 ↳ only the testcase inputs are
 ↳ required.

Program:
 {slow_code}

Explanation:
 {Your explanation here}

Test case Inputs:
 {Your testcase inputs}

Figure 11: Anchor Verification Framework Stage 1 (Test Inputs Generation) Prompt.

You are a code expert, and your
 ↳ task is to correct the
 ↳ functionally incorrect code
 ↳ based on test cases and
 ↳ execution feedback. Analyze the
 ↳ issues, apply the necessary
 ↳ fixes, and ensure the corrected
 ↳ code meets the expected
 ↳ functionality and pass the
 ↳ testcase.

Incorrect Program:
 {code}

Explanation:
 {explanation}

Testcase:
 {Testcase}

Feedback from execution:
 {Feedback}

Your corrected code version:

Figure 12: Anchor Verification Framework Stage 3 (Iterative Refinement) Prompt.

in Stage 3 of the Anchor Verification method, as depicted in Figure 12.

G Results of Anchor Verification Framework on PIE.

We conducted experiments using "QWEN2.5-CODER 32B fine-tuned on PIE" as the baseline and compared it with other methods. The results, shown in Table 4, demonstrate that Anchor Verification consistently delivers the high-


```
Below is a potentially problematic
↪ C++ program. Please provide a
↪ line-by-line explanation and
↪ correct any errors that may be
↪ present.
```

```
### Program:
{program}
```

```
### Explanation:
{Your explanation here}
```

```
### Revised Program:
{Your revised program here}
```

Figure 13: Self-Debugging Prompt.

```
Given the program below, please
↪ explain and analyze its
↪ functionality, and generate
↪ three comprehensive test cases
↪ that thoroughly cover boundary
↪ conditions and all code paths.
↪ Each testcase should include
↪ the input and the corresponding
↪ expected output.
```

```
### Program:
{slow_code}
```

```
### Explanation:
{Your explanation here}
```

```
### Test case:
{Your testcase}
```

Figure 14: The Prompt of Direct Test Generation Method.

H Detailed Examples of User-Oriented and Problem-Oriented Perspectives.

We provide detailed examples, as shown in Figure 16, Figure 17, and Figure 18, to illustrate that in the original PIE, program optimization pairs are constructed through iterative submissions and optimizations by the same user for the same programming problem, which can be limited by the single programmer’s thought patterns.

I Case Study of Anchor Verification Framework.

We present three case studies to vividly illustrate specific examples of Anchor Verification, as depicted in Figures 19, 20, and 21. These cases offer a clear and intuitive understanding of how Anchor Verification operates in practice.

est performance gains. On the DEEPSEEK-V3 backbone, we observed improvements in %OPT (31.24% → 47.28%), SPEEDUP (2.95× → 3.40×), and CORRECT (46.52% → 65.32%). Furthermore, we found that the gains in optimization ratio and speedup brought by the Anchor Verification Mechanism’s improvements in correctness for PIE were not as significant as those observed in PCO. For example, on the DEEPSEEK-V3 backbone, CORRECT increased by 18.8%, but SPEEDUP only improved by 0.45×. In contrast, in the PCO scenario, CORRECT increased by 12.99%, while SPEEDUP saw a larger improvement of 0.86×.

Table 4: Results of Anchor Verification and compared methods with QWEN2.5-CODER, GPT-4o, and DEEPSEEK-V3 on BEST@1. The improvement (denoted as Δ) is measured against the baseline (w/o refinement).

LLMs	Methods	BEST@1					
		%OPT	$\Delta \uparrow$	SPEEDUP	$\Delta \uparrow$	CORRECT	$\Delta \uparrow$
	Baseline (w/o refinement)	31.24%		2.95×		46.52%	
QWEN2.5-CODER 32B INSTRUCT	Self Debugging	35.69%	+4.45	3.02×	+0.07	53.74%	+7.22
	Direct Test Generation	38.74%	+7.50	3.08×	+0.13	57.49%	+10.97
	Anchor Verification	40.48%	+9.24	3.17×	+0.22	59.09%	+12.57
GPT-4o	Self Debugging	37.47%	+6.23	3.06×	+0.11	55.65%	+9.13
	Direct Test Generation	39.64%	+8.40	3.13×	+0.18	57.86%	+11.34
	Anchor Verification	42.50%	+11.26	3.32×	+0.37	63.60%	+17.08
DEEPSEEK-V3	Self Debugging	40.61%	+9.37	3.23×	+0.28	59.62%	+13.10
	Direct Test Generation	40.17%	+8.93	3.18×	+0.23	58.73%	+12.21
	Anchor Verification	47.28%	+16.04	3.40×	+0.45	65.32%	+18.80

```
#include <iostream>
#include <stdio.h>
using namespace std;
typedef long ll;

int main() {
    int length;
    ll arr[2000000];
    ll res[2000000] = {0};
    ll temp = 0;
    ll m = 2147483647;
    scanf("%d", &length);
    for (int i = 0; i < length; ++i) {
        scanf("%ld", &arr[i]);
    }
    res[0] = arr[0];
    for (int i = 1; i < length; ++i) {
        res[i] += res[i - 1] + arr[i];
    }
    for (int i = 1; i < length; ++i) {
        temp = abs(res[length - 1] - res[i - 1] * 2);
        m = min(temp, m);
    }
    printf("%ld\n", m);
    return 0;
}
```

(a) user1, initialization version.

```
#include <bits/stdc++.h>
using namespace std;

#define int long long
typedef vector<int> vi;

const int INF = 1e18 + 5;

void solve() {
    int n;
    cin >> n;
    vi v(n), pre(n);
    int mn = INF, s = 0;
    for (int i = 0; i < n; i++) cin >> v[i];
    pre[0] = v[0];
    for (int i = 1; i < n; i++) pre[i] = v[i] + pre[i - 1];
    for (int i = n - 1; i >= 1; i--) {
        s += v[i];
        mn = min(mn, abs(pre[i - 1] - s));
    }
    cout << mn;
}

signed main() {
    speed;
    int t = 1;
    while (t--) solve();
}
```

(b) user1, iteration version.

```
#include <cstdio>
const int MAX = 2e5 + 5;
int a[MAX];
int main() {
    int n;
    long long sum = 0;
    scanf("%d", &n);
    for (int i = 0; i < n; i++) {
        scanf("%d", &a[i]);
        sum += a[i];
    }
    long long left, right, temp;
    left = sum - a[n - 1];
    right = a[n - 1];
    long long min = left > right ? left - right : right - left;
    left = 0;
    for (int i = 0; i < n - 2; i++) {
        left += a[i];
        right = sum - left;
        temp = left > right ? left - right : right - left;
        if (temp < min) min = temp;
    }
    printf("%d\n", min);
    return 0;
}
```

(c) another user submitted version.

Figure 15: The three submitted code solutions all address problem "p03661", which asks for a split point in an array that minimizes the absolute difference between the sums of the two parts. Solutions (a) and (b) are different submissions from same user "u018679195". In (a), the prefix sum is calculated first, then the minimum difference is computed from start to finish. In (b), the prefix sum is also calculated first, but the minimum difference is computed from end to start, avoiding additional multiplication operations. Solution (c), from user "u353919145", calculates the difference between the left and right sums in real-time, requiring only one pass through the loop. It can be seen that solutions (a) and (b) only make local changes, while (c) constructs a more efficient algorithm.

```

#include <bits/stdc++.h>
using namespace std;

#define int long long

const int N = 1e5 + 5, M = 5, inf = 1e15;

int dp[N][M], a[N];

char op[N];

int Sign(int x) {
    if (x % 2) return -1;
    return 1;
}

int32_t main() {
    for (int i = 0; i < N; i++)
        for (int j = 0; j < M; j++) dp[i][j] = -inf;
    int n; cin >> n >> a[0];
    for (int i = 1; i < n; i++)
        cin >> op[i] >> a[i];
    dp[0][0] = a[0];
    for (int i = 1; i < n; i++)
        for (int j = M - 1; j >= 0; j--) {
            if (op[i] == '+') dp[i][j] =
                dp[i - 1][j] + a[i] * Sign(j);
            else if (j) dp[i][j] =
                dp[i - 1][j - 1] +
                a[i] * Sign(j);
            if (j + 1 < M) dp[i][j] =
                max(dp[i][j], dp[i][j + 1]);
        }
    cout << dp[n - 1][0] << "\n";
}

```

(a) user1, initialization version.

```

#include <bits/stdc++.h>
using namespace std;

#define int long long
const int N = 1e5 + 5, M = 3, inf = 1e15;

int dp[N][M], a[N];
char op[N];

int Sign(int x) {
    if (x % 2) return -1;
    return 1;
}

int32_t main() {
    ios::sync_with_stdio(0),
    cin.tie(0), cout.tie(0),
    cout.tie(0);
    for (int i = 0; i < N; i++)
        for (int j = 0; j < M; j++) dp[i][j] = -inf;
    int n; cin >> n >> a[0];
    for (int i = 1; i < n; i++)
        cin >> op[i] >> a[i];
    dp[0][0] = a[0];
    for (int i = 1; i < n; i++)
        for (int j = M - 1; j >= 0; j--) {
            if (op[i] == '+') dp[i][j] =
                dp[i - 1][j] + a[i] * Sign(j);
            else if (j) dp[i][j] =
                dp[i - 1][j - 1] +
                a[i] * Sign(j);
            if (j + 1 < M) dp[i][j] =
                max(dp[i][j], dp[i][j + 1]);
        }
    cout << dp[n - 1][0] << "\n";
}

```

(b) user1, iteration version.

```

#include <cstdio>
#include <algorithm>
using namespace std;
const int MAXN = 1e5 + 5;
typedef long long LL;
#define INF LL(1e15)
LL s1, s2, as, n;
LL sz[MAXN], fh[MAXN];
char c[5];
int main()
{
    scanf("%lld", &n);
    scanf("%lld", &as);
    getchar();
    for (LL i = 1; i <= n - 1; i++) {
        scanf("%s", c);
        scanf("%d", &sz[i]);
        fh[i] = c[0];
    }
    s1 = s2 = -INF;
    for (LL i = 1; i <= n - 1; i++) {
        if (fh[i] == '-') {
            as -= sz[i];
            s1 -= sz[i];
            s2 += sz[i];
            s1 = max(s1, s2);
            s2 = max(as, s2);
        }
        else {
            as += sz[i];
            s1 += sz[i];
            s2 -= sz[i];
        }
        s2 = max(s1, s2);
        as = max(s2, as);
    }
    printf("%lld", as);
}

```

(c) another user submitted version.

Figure 16: The above three code snippets all come from the problem "p03580", which involves maximizing the evaluated value of a given formula by adding an arbitrary number of pairs of parentheses and outputting the maximum possible value. (a) and (b) are from the same user "u1821171064", both employing dynamic programming algorithms with a time complexity of $\mathcal{O}(N * M)$, where N is the length of the sequence and M is the number of states. In (b), the number of states M is reduced, and input and output are optimized. (c) is from user "u863370423" and uses a greedy algorithm, which is suitable for problems with fewer current states where the global optimal solution can be achieved through local optimization, with a time complexity of $\mathcal{O}(N)$.

```

#include <iostream>
#include <cstring>
using namespace std;
typedef long long LL;
#define F(i) for(int i=0;i<n;i++)

int d[555][555] = {0}, c[555][555]
    = {0};

int qu(int l, int r) {
    if (l > r) return 0;
    if (d[l][r] != -1) return
        d[l][r];
    return d[l][r] = c[l][r] +
        qu(l + 1, r) + qu(l, r -
            1) - qu(l + 1, r - 1);
}

int main() {
    memset(d, -1, sizeof(d));
    int n, m, q;
    cin >> n >> m >> q;
    while (m--) {
        int l, r;
        cin >> l >> r;
        c[l][r]++;
    }
    while (q--) {
        int l, r;
        cin >> l >> r;
        cout << qu(l, r) << endl;
    }
    return 0;
}

```

(a) user1, initialization version.

```

#include <bits/stdc++.h>
using namespace std;

#define int long long
#define pb push_back
#define faster
    ios::sync_with_stdio(0)

const int N = 509;
vector<int> v[N + 5];

int32_t main() {
    faster;
    int n, p, q;
    cin >> n >> p >> q;
    int x, y;
    for (int i = 1; i <= p; i++) {
        cin >> x >> y;
        v[x].pb(y);
    }
    for (int i = 1; i <= n; i++) {
        sort(v[i].begin(),
            v[i].end());
    }
    while (q--) {
        cin >> x >> y;
        int ans = 0;
        for (int i = x; i <= y;
            i++) {
            ans += upper_bound(
                v[i].begin(),
                v[i].end(), y)
                - v[i].begin();
        }
        cout << ans << "\n";
    }
    return 0;
}

```

(b) user1, iteration version.

```

#include <cstdio>
#define int long long
#define dotimes(i, n) for (int i =
    0; i < (n); i++)

using namespace std;

int rint() {
    int n;
    scanf("%lld", &n);
    return n;
}

void wint(int n) {
    printf("%lld\n", n);
}

signed main() {
    int N = rint();
    int M = rint();
    int Q = rint();
    int S[N + 1][N + 1];
    dotimes(R, N + 1)
        dotimes(L, N + 1)
            S[R][L] = 0;
    dotimes(i, M) {
        int L = rint();
        int R = rint();
        S[R][L]++;
    }
    dotimes(R, N)
        dotimes(L, N)
            S[R + 1][L + 1] += S[R +
                1][L] + S[R][L + 1] -
                S[R][L];
    dotimes(i, Q) {
        int p = rint() - 1;
        int q = rint();
        wint(S[q][q] + S[p][p] -
            S[q][p] - S[p][q]);
    }
    return 0;
}

```

(c) another user submitted version.

Figure 17: The above three code segments all come from the same problem "p03283", which deals with cumulative sum queries in a 2D matrix. (a) and (b) are different submission versions from the same user "u816631826". In (a), the problem is solved using recursion and dynamic programming, but the query time complexity is high, $\mathcal{O}(N^2)$. In (b), the STL-provided binary search function is used, reducing the time complexity to $\mathcal{O}(N * \log(N))$. (c) comes from another user "u281670674" and solves the problem using a 2D prefix sum matrix. The preprocessing time complexity is $\mathcal{O}(N^2)$, but the query time complexity for each query is $\mathcal{O}(1)$, making it more efficient.


```
#include <bits/stdc++.h>

using namespace std;

inline void rd(int &x) {
    char ch;
    for (; !isdigit(ch=getchar()););
    for (x=ch-'0';
    isdigit(ch=getchar());)
        x=x*10+ch-'0';
}

typedef long long LL;

const int MAXN = 300005;

int N, n, a[MAXN], cnt[MAXN];

LL sum[MAXN];

int ans[MAXN];

inline bool chk(int k, int x) {
    int pos = upper_bound(a + 1, a + n + 1, x) - a;
    return sum[pos-1] +
        1ll*(n-pos+1)*x >=
        1ll*k*x;
}

int main() {
    rd(N);
    for (int i = 1, x; i <= N; ++i)
        rd(x), ++cnt[x];
    for (int i = 1; i <= 300000;
        ++i) if (cnt[i]) a[++n] =
        cnt[i];
    sort(a + 1, a + n + 1);
    for (int i = 1; i <= n; ++i)
        sum[i] = sum[i-1] + a[i];
    int now = 0;
    for (int k = n; k >= 1; --k) {
        while (now < N && chk(k,
            now+1)) ++now;
        ans[k] = now;
    }
    for (int i = 1; i <= N; ++i)
        printf("%d\n", ans[i]);
}
```

(a) user1, initialization version.

```
#include <bits/stdc++.h>

using namespace std;

inline void rd(int &x) {
    char ch;
    for (; !isdigit(ch=getchar()););
    for (x=ch-'0';
    isdigit(ch=getchar());)
        x=x*10+ch-'0';
}

typedef long long LL;

const int MAXN = 300005;

int n, cnt[MAXN];

LL sum[MAXN];

int ans[MAXN];

inline bool chk(int k, int x) {
    return sum[x] >= 1ll*k*x;
}

int main() {
    rd(n);
    for (int i = 1, x; i <= n; ++i)
        rd(x), ++cnt[x],
        ++sum[cnt[x]];
    for (int i = 1; i <= n; ++i)
        sum[i] += sum[i-1];
    int now = 0;
    for (int k = n; k >= 1; --k) {
        while (now < n && chk(k,
            now+1)) ++now;
        ans[k] = now;
    }
    for (int i = 1; i <= n; ++i)
        printf("%d\n", ans[i]);
}
```

(b) user1, iteration version.

```
#include <bits/stdc++.h>
#include <cstdio>
using namespace std;
typedef long long ll;
#define rep(i, n) for(int i = 0; i
< (n); i++)
#define repl(i, n) for(int i = 1;
i <= (n); i++)

int hist[300002], cnt[300001];
const int cm = 1 << 17;
char cn[cm], * ci = cn + cm, ct;

inline char getcha() {
    if (ci - cn == cm) {
        fread_unlocked(cn, 1, cm,
        stdin); ci = cn; }
    return *ci++;
}
inline int getint() {
    int A = 0;
    if (ci - cn + 16 > cm) while
((ct = getcha()) >= '0') A
= A * 10 + ct - '0';
    else while ((ct = *ci++) >=
'0') A = A * 10 + ct -
'0';
    return A;
}

const int dm = 1 << 21;
char dn[dm], * di = dn;
inline void putint(int X) {
    int keta = 0;
    char C[10];
    while (X) {
        *(C + keta) = '0' + X %
        10;
        X /= 10;
        keta++;
    }
    for (int i = keta - 1; i >= 0;
        i--) * di++ = (*C + i);
    *di++ = '\n';
}

int main() {
    int N = getint();
    rep(i, N) hist[getint()]++;
    repl(i, N) cnt[hist[i]]++;
    int k = 1;
    rep(i, N + 1) rep(j, cnt[i])
        hist[k++] = i

    k = N + 1;
    int ruiseki = N;
    int mae = 0;
    for (int i = N; i >= 1; i--) {
        while (hist[k - 1] >= i) {
            ruiseki -= hist[--k];
        }
        int kei = N - k + 1 +
            ruiseki / i;

        for (int j = mae + 1; j <=
            kei; j++) putint(i);
        mae = kei;
    }
    for (int j = mae + 1; j <= N;
        j++) {
        *di++ = '0';
        *di++ = '\n';
    }
    fwrite(dn, 1, di - dn,
        stdout);
    return 0;
}
```

(c) another user submitted version.

Figure 18: The above three code snippets all come from the problem "p02890", which requires calculating, for each possible K value (from 1 to N), the maximum number of times K cards with different numbers can be selected and removed from N cards. (a) and (b) are from the same user "u990400947" and utilize prefix sum calculation and searching. The latter employs condition checking with a time complexity of $\mathcal{O}(N * \log(N))$. (c) uses a difference array, reducing the time complexity to $\mathcal{O}(N)$.

```

** Optimized but incorrect Code **

#include <stdio.h>
#include <algorithm>
using namespace std;
int n,k,i,a[100005];
int main()
{
    scanf("%d%d",&n,&k);
    for(i=1;i<=n;i++) scanf("%d",&a[i]);
    sort(a+1,a+n+1);
    printf("%d\n",n-lower_bound(a+1,a+n+1,k)+1);
    return 0;
}

```

Below is the Explanation of slow code and generated test case inputs

****Explanation:****

This code counts how many elements in an array are greater than or equal to a
 ↪ given value k. Here is the breakdown:

Input: Reads n (array size) and k (threshold value), followed by an array h
 ↪ of size n.

Sorting: Sorts the array h in descending order.

Counting: Iterates through the sorted array and counts how many elements are
 ↪ greater than or equal to k.

Output: Prints the count.

****Test Inputs:****

Test case 1 input:

5 10
 15 8 12 10 20

Test case 2 input:

4 5
 1 2 3 4

Test case 3 input:

6 7
 7 7 7 7 7 7

**** Verified Test cases ****

**** Test Cases****

Test case 1 input:

5 10
 15 8 12 10 20

Test case 1 output:

4

Test case 2 input:

4 5
 1 2 3 4

Test case 2 output:

0

Test case 3 input:

6 7
 7 7 7 7 7 7

Test case 3 output:

6

**** Optimized and Correct Code****

```

#include <stdio.h>
#include <algorithm>
using namespace std;
int n, k, i, a[100005];
int main() {
    scanf("%d%d", &n, &k);
    for (i = 0; i < n; i++)
        ↪ scanf("%d", &a[i]);
    sort(a, a + n);
    printf("%d\n", n -
        ↪ (lower_bound(a, a + n, k) -
        ↪ a));
    return 0;
}

```

Figure 19: Case Study 1. The reason for the error is that the offset was not calculated correctly. The correct fix is to calculate the index using `lower_bound(...)` - `a`.

**** Optimized but incorrect Code ****

```
#include <stdio>
#include <algorithm>
using std::min;
#define MAXN 100003
int abs(int z)
{
    if (z<0) return -z;
    else return z;
}
int main(void)
{
    int N,K,i,j;
    int h[MAXN], p[MAXN];
    scanf("%d%d",&N,&K);
    for (i=0;i<N;i++)
    {
        scanf("%d",h+i);
    }
    p[N-1]=0;
    p[N-2]=abs(h[N-2]-h[N-1]);
    for (i=N-3;i>=0;i--)
    {
        p[i]=min(p[i+1]+abs(h[i]-h[i+1]),p[i+2]
            +abs(h[i]-h[i+2]));
    }
    printf("%d\n",p[0]);
    return 0;
}
```

*Below is the Explanation of slow code and generated test **case inputs***

****Explanation:****

This code solves a dynamic programming problem to find the minimum cost to
→ reach the n-th element from the first element in an array, where you can
→ jump up to k steps back at each position. Here's the breakdown:

Input: Reads n (array size) and k (maximum jump distance), followed by the
→ array num.

Initialization:

dp array is initialized to a large value, except dp[0] and dp[1], which are
→ set to 0 (starting points with no cost).

DP Transition:

For each position i from 2 to n, it calculates the minimum cost to reach i by
→ considering jumps from positions j (within [i-k, i-1]).

The cost is the absolute difference between num[i] and num[j], added to the
→ cost to reach j.

Output: The minimum cost to reach the n-th element (dp[n]) is printed.

**** Test Cases****

Test **case 1 input:**

```
5 2
1 3 5 2 4
```

Test **case 2 input:**

```
4 1
10 20 10 30
```

Test **case 3 input:**

```
6 3
1 2 3 4 5 6
```

```
** Verified Test cases **
```

```
** Test Cases**
```

```
Test case 1 input:
```

```
5 2
```

```
1 3 5 2 4
```

```
Test case 1 output:
```

```
5
```

```
Test case 2 input:
```

```
4 1
```

```
10 20 10 30
```

```
Test case 2 output:
```

```
40
```

```
Test case 3 input:
```

```
6 3
```

```
1 2 3 4 5 6
```

```
Test case 3 output:
```

```
5
```

```
** Optimized and Correct Code**
```

```
#include <stdio>
```

```
#include <algorithm>
```

```
#include <climits>
```

```
using std::min;
```

```
#define MAXN 100003
```

```
int abs(int z)
```

```
{
```

```
    if (z<0)
```

```
        return -z;
```

```
    else
```

```
        return z;
```

```
}
```

```
int main(void)
```

```
{
```

```
    int N,K,i,j;
```

```
    int h[MAXN], p[MAXN];
```

```
    scanf("%d%d",&N,&K);
```

```
    for (i=0;i<N;i++)
```

```
    {
```

```
        scanf("%d",h+i);
```

```
    }
```

```
    p[0] = 0;
```

```
    for (i=1;i<N;i++)
```

```
    {
```

```
        p[i] = INT_MAX;
```

```
        for (j=1;j<=K && i-j>=0;j++)
```

```
        {
```

```
            p[i] = min(p[i], p[i-j] + abs(h[i] - h[i-j]));
```

```
        }
```

```
    }
```

```
    printf("%d\n",p[N-1]);
```

```
    return 0;
```

```
}
```

Figure 20: Case Study 2. The reason for the error is that the code does not properly handle the jump limit K , and the state transition logic is incomplete, only considering the cost of adjacent positions and ignoring larger jump ranges. The correct modification is to use a double loop to dynamically update the minimum cost for each position by considering jumps up to K steps backward, thus fully addressing all possible jump scenarios.


```

** Optimized but incorrect Code **

#include<stdio>
#include<cstring>
#include<algorithm>
using namespace std;
int n,ans;
int a[10010];

int main(){
    scanf("%d",&n);
    for(int i=1;i<=n;i++) scanf("%d",&a[i]);
    while(1){
        int mx=a[1],id=1;
        for(int i=2;i<=n;i++)
            if(a[i]>mx) mx=a[i],id=i;
        if(mx<n-1) break;
        ans+=mx/n;
        a[id]%=n;
        for(int i=1;i<=n;i++)
            if(i!=id) a[i]+=mx/n;
    }
    printf("%d\n",ans);
}

```

Below is the Explanation of slow code and generated test case inputs

****Explanation:****

The code simulates a process where an array a of size n is repeatedly
 ↳ modified until the maximum value in the array is no greater than n-1. The
 ↳ process involves the following steps:

Input: The code reads the size of the array n and the elements of the array
 ↳ a.

Main Loop: The code enters an infinite loop that continues until the maximum
 ↳ value in the array is less than or equal to n-1. In each iteration:
 The maximum value maxn in the array and its index id are identified.
 If maxn is less than or equal to n-1, the loop terminates.

Otherwise:

The value at a[id] is updated to maxn % n (the remainder when maxn is divided
 ↳ by n).

All other elements in the array are incremented by maxn / n (the quotient
 ↳ when maxn is divided by n).

The operation count ans is incremented by maxn / n.

Output: The total number of operations ans is printed.

**** Test Cases****

Test case 1 input:

2
0 7

Test case 2 input:

4
8 0 0 0

Test case 3 input:

4
1 8 3 4

```
** Verified Test cases **
```

```
** Test Cases**
```

```
Test case 1 input:
```

```
2
```

```
0 7
```

```
Test case 1 output:
```

```
6
```

```
Test case 2 input:
```

```
4
```

```
8 0 0 0
```

```
Test case 2 output:
```

```
2
```

```
Test case 3 input:
```

```
4
```

```
1 8 3 4
```

```
Test case 3 output:
```

```
9
```

```
** Optimized and Correct Code**
```

```
#include<stdio>
```

```
#include<cstring>
```

```
#include<algorithm>
```

```
using namespace std;
```

```
int n,ans;
```

```
int a[10010];
```

```
int main() {
```

```
    scanf("%d",&n);
```

```
    for(int i=1;i<=n;i++) scanf("%d",&a[i]);
```

```
    while(1) {
```

```
        int mx=a[1],id=1;
```

```
        for(int i=2;i<=n;i++)
```

```
            if(a[i]>mx) mx=a[i],id=i;
```

```
        if(mx<n) break;
```

```
        ans+=mx/n;
```

```
        a[id]%=n;
```

```
        for(int i=1;i<=n;i++)
```

```
            if(i!=id) a[i]+=mx/n;
```

```
    }
```

```
    printf("%d\n",ans);
```

```
    return 0;
```

```
}
```

Figure 21: Case Study 3. The error is that the termination condition used $mx < n - 1$, which prematurely stopped the loop, while the correct condition is $mx < n$, ensuring the loop only stops when maximum value is less than n .